



Free Questions for AZ-104

Shared by Burgess on 06-03-2026

For More Free Questions and Preparation Resources

[Check the Links on Last Page](#)



Question 1

Question Type: MultipleChoice

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an app named App1 that is installed on two Azure virtual machines named VM1 and VM2. Connections to App1 are managed by using an Azure Load Balancer.

The effective network security configurations for VM2 are shown in the following exhibit.



You discover that connections to App1 from 131.107.100.50 over TCP port 443 fail. You verify that the Load Balancer rules are configured correctly.

You need to ensure that connections to App1 can be established successfully from 131.107.100.50 over TCP port 443.

Solution: You create an inbound security rule that denies all traffic from the 131.107.100.50 source and has a cost of 64999. 

Does this meet the goal?

Options:

A- Yes

B- No

Answer:

B

Explanation:

In the provided NSG configuration for VM2, the effective inbound security rules show the following priorities:

Priority

Name

Port

Protocol

Source

Action

100

Allow_131.107.100.50

443

TCP

131.107.100.50

Allow

200

BlockAllOther443

443

TCP



Any

Deny

65000+

Default rules

Various

Any

Allow/Deny

The problem states that connections from 131.107.100.50 to port 443 fail, even though the Load Balancer rules are correct.

The proposed solution is to create a new inbound rule that denies all traffic from 131.107.100.50 with a priority of 64999. However, this does not solve the problem --- it makes it worse, because:

The rule with priority 64999 has a higher number, meaning it is evaluated later than existing rules.

Since NSG rules are processed in ascending order of priority, the first matching rule applies.

The rule at priority 200 (BlockAllOther443) already denies the connection before the new rule is even evaluated.

To fix the issue, the correct solution would be to modify or remove the "BlockAllOther443" rule, or adjust priorities so that the allow rule for 131.107.100.50 is evaluated first.

Simply adding another deny rule with a lower priority number (higher numeric value) will not override existing denies.

Final Verified Answer: B. No

Question 2

Question Type: MultipleChoice

You have an Azure subscription that contains the resources shown in the following table.



You need to configure a proximity placement group for VMSS1.

Which proximity placement groups should you use?



Options:

- A- Proximity2 only
- B- Proximity 1, Proximity2, and Proximity3
- C- Proximity 1 and Proximity3 only
- D- Proximity1 only

Answer:

A

Explanation:

Placement Groups is a capability to achieve co-location of your Azure Infrastructure as a Service (IaaS) resources and low network latency among them, for improved application performance.

Azure proximity placement groups represent a new logical grouping capability for your Azure Virtual Machines, which in turn is used as a deployment constraint when selecting where to place your virtual machines. In fact, when you assign your virtual machines to a proximity placement group, the virtual machines are placed in the same data center, resulting in lower and

deterministic latency for your applications.

The VMSS should share the same region, even it should be the same zone as proximity groups are located in the same data center. Accordingly, it should be proximity 2 only.

<https://azure.microsoft.com/en-us/blog/introducing-proximity-placement-groups>

Question 3

Question Type: MultipleChoice

You have an Azure subscription named Subscription1 that contains virtual network named VNet1. VNet1 is in a resource group named RG1. A user named User1 has the following roles for Subscription1:

- * Reader
- * Security Admin
- * Security Reader

You need to ensure that User1 can assign the Reader role for VNet1 to other users. What should you do?



Options:

- A- Remove User1 from the Security Reader and Reader roles for Subscription1.
- B- Assign User1 the Owner role for VNet1.
- C- Remove User1 from the Security Reader role for Subscription1. Assign User1 the Contributor role for RG1.
- D- Remove User1 from the Security Reader and Reader roles for Subscription1. Assign User1 the Contributor role for Subscription1

Answer:

B

Explanation:

<https://docs.microsoft.com/en-us/azure/role-based-access-control/rbac-and-directory-admin-roles#:~:text=The%20User%20Access%20Administrator%20role%20enables%20the%20user%20to%20grant,Azure%20subscriptions%20and%20management%20groups.>

Question 4

Question Type: MultipleChoice

You have an Azure Storage account named storage1.

You need to enable a user named User1 to list and regenerate storage account keys for storage1.

Solution: You assign the Storage Account Encryption Scope Contributor Role to User1.

Does this meet the goal?

Options:

A- Yes

B- No

Answer:

B



Explanation:

In Azure, storage account keys are managed through Azure Role-Based Access Control (RBAC) permissions. The ability to list and regenerate storage account keys requires access to specific management operations on the storage account resource, not encryption scopes.

According to Microsoft Azure Administrator documentation, the permissions required to list and regenerate storage account keys are:

Microsoft.Storage/storageAccounts/listkeys/action

Microsoft.Storage/storageAccounts/regeneratekey/action

Only specific built-in roles include these actions, such as:

Owner

Contributor

Storage Account Contributor

However, the role mentioned in the question --- Storage Account Encryption Scope Contributor --- is designed for a different purpose.

From Microsoft's official documentation for Storage Account Encryption Scope Contributor:

"Grants permissions to manage encryption scopes within a storage account. Does not grant access to manage storage account keys or data."

This means that assigning the Storage Account Encryption Scope Contributor role gives the user permissions to configure encryption scopes (for example, enabling or modifying encryption at the blob container level), but not to list or regenerate access keys.

Therefore, the solution fails to meet the requirement of enabling User1 to list and regenerate keys.

To achieve the goal, the correct role would be: Storage Account Contributor

This role explicitly allows managing all storage account properties, including key management.

Key Point (Microsoft Official Documentation Summary):

Storage Account Encryption Scope Contributor Manage encryption scopes only.

Storage Account Contributor Manage storage account settings, including listing and regenerating keys.

Question 5



Question Type: MultipleChoice

You plan to create the Azure web apps shown in the following Table.



What is the minimum number of App Service plans you should create for the web apps?

Options:

- A- 1
- B- 2
- C- 3
- D- 4

Answer:

B

Explanation:

NET Core 3.0: Windows and Linux ASP .NET V4.7: Windows only PHP 7.3: Windows and Linux Ruby 2.6: Linux only Also, you can't use Windows and Linux Apps in the same App Service Plan, because when you create a new App Service plan you have to choose the OS type. You can't mix Windows and Linux apps in the same App Service plan. So, you need 2 ASPs. Reference: <https://docs.microsoft.com/en-us/azure/app-service/overview>

Question 6

Question Type: MultipleChoice

You have an Azure web app named App1. App1 has the deployment slots shown in the following table:



In webapp1-test, you test several changes to App1.

You back up App1.

You swap webapp1-test for webapp1-prod and discover that App1 is experiencing performance issues.

You need to revert to the previous version of App1 as quickly as possible.

What should you do?

Options:

- A- Redeploy App1
- B- Swap the slots
- C- Clone App1
- D- Restore the backup of App1

Answer:

B

Explanation:

When you swap deployment slots, Azure swaps the Virtual IP addresses of the source and destination slots, thereby swapping the URLs of the slots. We can easily revert the deployment by

swapping back. Deployment slots are live apps with their own host names. App content and configurations elements can be swapped between two deployment slots, including the production slot. Deploying your application to a non-production slot has the following benefits: 1. You can validate app changes in a staging deployment slot before swapping it with the production slot. 2. Deploying an app to a slot first and swapping it into production makes sure that all instances of the slot are warmed up before being swapped into production. Reference: <https://docs.microsoft.com/en-us/azure/app-service/deploy-staging-slots>

Question 7

Question Type: MultipleChoice

You have an Azure subscription.

Users access the resources in the subscription from either home or from customer sites. From home, users must establish a point-to-site VPN to access the Azure resources. The users on the customer sites access the Azure resources by using site-to-site VPNs.

You have a line-of-business app named App1 that runs on several Azure virtual machine. The virtual machines run Windows Server 2016.

You need to ensure that the connections to App1 are spread across all the virtual machines.

What are two possible Azure services that you can use? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

Options:

- A- a public load balancer
- B- Traffic Manager
- C- an Azure Content Delivery Network (CDN)
- D- an internal load balancer
- E- an Azure Application Gateway

Answer:

D, E

Explanation:

Line of Business WebAPP works on VMs need internal load balancer. So D is needed. Then deploy WebAPP on VMs, check the link.
<https://docs.microsoft.com/en-us/azure/application-gateway/quick-create-portal> So B is needed as well. The original answer is not accomplished.

Question 8

Question Type: MultipleChoice

You need to implement the planned changes for the storage account content. Which containers and file shares can you use to organize the content?

Options:

- A- share1 only
- B- cont1 and share1 only
- C- share1 and share2 only
- D- cont1, share1, and share2 only
- E- cont1, cont2, share1, and share2



Answer:

B

Explanation:

In the scenario, storage1 is configured as StorageV2 with Hierarchical namespace = Yes, while storage2 is configured as StorageV2 with Hierarchical namespace = No.

From Microsoft's Azure Storage Documentation and AZ-104 Study Guide, the following principles apply:

A hierarchical namespace (enabled when the storage account has Azure Data Lake Storage Gen2 capabilities) allows the use of directories within containers to organize data.

The hierarchical namespace provides directory and file-level structure similar to a file system. This is supported only for blob containers, not for Azure Files.

Azure Files (file shares) do not depend on hierarchical namespaces and cannot have directories in the same way Data Lake Gen2 does --- directories can exist inside the share but not in the blob container sense.

The planned change states that you must use directories whenever possible to organize content. Therefore, only storage accounts with hierarchical namespace enabled can use directory structures --- that's storage1.

In this case:

storage1 (Hierarchical namespace = Yes) supports containers (like cont1) and file shares (like share1).

storage2 (Hierarchical namespace = No) does not support directories within blob containers (Data Lake structure).

Hence, you can use only cont1 (container in storage1) and share1 (file share in storage1) to organize content as required.

This is directly supported by the Microsoft documentation on Data Lake Storage Gen2:

"When you enable the hierarchical namespace for a storage account, you can organize objects into directories and subdirectories. This capability is available only for accounts configured for Data Lake Storage Gen2."

Final Verified Answer: B. cont1 and share1 only

Question 9



Question Type: MultipleChoice

You need to configure an Azure web app named contoso.azurewebsites.net to host www.contoso.com.

What should you do first?

Options:

- A- Create a CNAME record named asuid that contains the domain verification ID.
- B- Create A records named www.contoso.com and asuid.contoso.com.
- C- Create a TXT record named asuid that contains the domain verification ID.
- D- Create a TXT record named www.contoso.com that has a value of contoso.azurewebsites.net.

Answer:

A

Explanation:

When you configure a custom domain (like `www.contoso.com`) for an Azure Web App (App Service), Azure requires that the domain be verified to ensure ownership before binding it to the app.

According to Microsoft Azure official documentation ("Map a custom domain name to your Azure web app" --- Microsoft Learn):

"Before you can add a custom domain, you must verify that you own the domain name by creating a DNS record with your domain registrar. Azure uses an asuid verification record, which can be either a CNAME or TXT record, depending on your DNS provider."

The first step is to create a CNAME or TXT record in your DNS zone that links your custom domain to the Azure verification ID.

The record name is `asuid.contoso.com`.

The value (target) is the domain verification ID shown in the Azure portal under Custom domains Custom hostnames Domain ownership.

Once Azure verifies the domain ownership using that record, you can add `www.contoso.com` as a custom hostname in the web app configuration.

Key Point:



Step 1: Verify domain ownership using a CNAME or TXT record.

Step 2: Bind the custom hostname (`www.contoso.com`) to your web app.

Thus, the correct and verified answer is: A. Create a CNAME record named `asuid` that contains the domain verification ID.

Question 10

Question Type: MultipleChoice

You develop the following Azure Resource Manager (ARM) template to create a resource group and deploy an Azure Storage account to the resource group.

Which cmdlet should you run to deploy the template?

Options:

A- `New-AzTenantDeployment`

- B- New-AzResourceGroupDeployment
- C- New-AzResource
- D- New-AzOeployment

Answer:

B

Explanation:

In Microsoft Azure, Azure Resource Manager (ARM) templates are JSON-based files used to define and automate the deployment of Azure resources in a consistent and repeatable manner. To deploy an ARM template, the PowerShell cmdlet you use depends on the deployment scope --- whether the template targets a tenant, management group, subscription, or resource group.

According to the Microsoft Azure Administrator documentation (AZ-104 Study Guide, "Deploy and Manage Azure Resource Manager Templates"):

New-AzTenantDeployment Used for tenant-level deployments, typically involving management group creation or tenant-wide policies.

New-AzSubscriptionDeployment Used for subscription-level deployments, when resources are created directly under the subscription rather than within a resource group.



New-AzResourceGroupDeployment Used for resource group-level deployments, which is the most common scenario for creating and managing resources such as virtual machines, storage accounts, and virtual networks.

New-AzDeployment An older alias that may refer to subscription-level deployments in older Az modules but is not the standard for resource group deployments.

In this question, the ARM template defines both a resource group and a storage account that will be deployed within that resource group. Therefore, the correct PowerShell cmdlet to use is:

`New-AzResourceGroupDeployment ``

`-ResourceGroupName <RGName> ``

`-TemplateFile <PathToTemplateFile>`

This cmdlet processes the specified ARM template file, validates the schema, and provisions all the resources defined in the template inside the target resource group.

Hence, based on the Azure Resource Manager official documentation and exam objectives, the correct answer is:

Question 11

Question Type: MultipleChoice

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have a Microsoft Entra tenant named contoso.com.

You have a CSV file that contains the names and email addresses of 500 external users.

You need to create a guest user account in contoso.com for each of the 500 external users.

Solution; From Microsoft Entra ID in the Azure portal, you use the Bulk invite users' operation.

Does this meet the goal?

Options:

A- Yes

B- No

Answer:

A

Explanation:

This scenario involves inviting 500 external users (B2B guests) into your Microsoft Entra tenant (formerly Azure Active Directory). The goal is to create guest user accounts efficiently using a CSV file containing user details (names and email addresses).

Understanding the Requirement

You must onboard external users (not internal directory users). These external accounts require B2B guest invitations, not normal user creation.

There are two supported approaches for bulk guest invitations:

Using the Microsoft Entra portal (Bulk invite users)

Using PowerShell with the New-AzureADMSInvitation cmdlet.

The Bulk invite users feature in the Entra portal is specifically designed for this exact use case --- bulk creation of guest accounts using a CSV file.

How the Bulk Invite Process Works

When you go to Microsoft Entra ID Users Bulk operations Bulk invite, you can:

Upload a CSV file containing columns like Email address, Display name, and First/Last names.

The portal validates the file format and initiates invitations for all 500 external users.

Each external user receives an invitation email and is added as a guest (UserType = Guest) in the directory.

Once accepted, these users can be assigned roles, access resources, or be part of Azure AD groups and applications.

Why This Meets the Goal

The solution correctly leverages the Bulk invite users functionality designed for mass external user onboarding. It satisfies all the requirements:

Creates guest accounts (not internal users).

Uses a CSV file, allowing batch import.

Achieves the goal using Microsoft Entra ID portal,  without needing PowerShell scripting.

Reference from Microsoft Documentation (Microsoft Learn -- Verified Source)

"You can use the Azure portal to invite multiple external users to your organization in one go using the Bulk invite feature. This process imports a CSV file containing user details and sends invitations to all users automatically." (Source: Microsoft Learn --- Bulk invite B2B users to Microsoft Entra ID)

Question 12

Question Type: MultipleChoice

You have an Azure App Service app named App1 that contains two running instances.

You have an autoscale rule configured as shown in the following exhibit.



For the Instance limits scale condition setting, you set Maximum to 5.

During a 30-minute period, App1 uses 80 percent of the available memory.

What is the maximum number of instances for App1 during the 30-minute period?

Options:

- A- 2
- B- 3
- C- 4
- D- 5

Answer:

D

Explanation:

In Azure App Service, autoscale rules automatically adjust the number of running instances based on performance metrics such as CPU, memory, or custom metrics. The configuration shown uses the Memory Percentage metric with a threshold of 70%, meaning that if average memory utilization exceeds 70% for 15 minutes, Azure will trigger a scale-out action.

From the exhibit, the rule specifies:



Metric threshold: Greater than 70%

Duration: 15 minutes

Action: Increase count by 1 instance

Cool down: 5 minutes

Maximum instance limit: 5

App1 currently has 2 running instances. If App1 maintains 80% memory utilization for 30 minutes, the autoscale mechanism will trigger the scale-out action after each qualifying 15-minute window, adding 1 instance every 5-minute cooldown period until the maximum instance limit (5) is reached.

Therefore, the progression would be:

Start: 2 instances

After first 15 minutes: +1 instance 3 total

After next 5-minute cooldown and continuing high memory +1 instance 4 total

After another cycle +1 instance 5 total

Since the rule continues to trigger until the maximum (5) is reached, and the memory usage remains above the threshold, the maximum number of instances that App1 will scale to during the 30-minute period is 5.



To Get Premium Files for AZ-104 Visit

<https://www.p2pexams.com/products/az-104>

For More Free Questions Visit

<https://www.p2pexams.com/microsoft/pdf/az-104>

20%
DISCOUNT

